

Intrusion-Detection-Model

1 Mounting Drive

```
[ ]: from google.colab import drive  
drive.mount('/content/drive')
```

2 Installations

```
[ ]: !pip install tensorflow --upgrade
```

```
[ ]: !pip install keras --upgrade
```

```
[ ]: !pip install yellowbrick==1.0.1
```

3 Imports

```
[ ]: import torch  
import torch.nn as nn  
from torch.utils.data import Dataset, DataLoader  
import os  
import numpy as np  
import pandas as pd  
import torch.optim as optim  
import torch.nn.functional as F  
import matplotlib.pyplot as plt  
from sklearn.decomposition import PCA  
import sklearn.cluster as cluster  
from mpl_toolkits.mplot3d import axes3d  
# from random import sample, shuffle  
import joblib  
from yellowbrick.cluster import KElbowVisualizer  
from scipy.stats import zscore, multivariate_normal  
import seaborn as sns
```

```

from sklearn.metrics import confusion_matrix, f1_score, make_scorer, \
    accuracy_score, precision_score, recall_score, \
    precision_recall_curve, roc_curve, auc
from sklearn.model_selection import GridSearchCV
from sklearn.ensemble import RandomForestClassifier, AdaBoostClassifier
from sklearn.utils import shuffle
from keras import Sequential
from keras.layers import Dense, Dropout, Convolution1D, MaxPooling1D, LSTM
from keras.utils import to_categorical
from keras.regularizers import l2
from keras.optimizers import Adam
from keras.utils import plot_model
from imblearn.under_sampling import RandomUnderSampler
from imblearn.over_sampling import SMOTE, SMOTENC, SVMSMOTE
from collections import Counter

```

```

[ ]: import warnings
     warnings.filterwarnings('ignore')

```

4 Constants

```

[ ]: CSV_DIR = '/content/drive/My Drive/Colab Notebooks/Intrusion Detection/data/
     ↳NSL-KDD/csv'
     BATCH_SIZE = 128
     FIG_TITLE_SIZE = 15
     LEGEND_TEXT_SIZE = 11
     MODELS_DIR = '/content/drive/My Drive/Colab Notebooks/Intrusion Detection/
     ↳saved_models/grid_search'

```

5 Device Check for GPU

```

[ ]: device = torch.device('cuda: 0' if torch.cuda.is_available() else 'cpu')
     print(device)

```

6 Dataset & DataLoader

6.1 NSLKDDDataset

```
[ ]: class NSLKDDDataset(Dataset):

    def __init__(self, features, labels):
        self.features = np.asarray(features)
        self.labels = np.asarray(labels)

    def __len__(self):
        return self.features.shape[0]

    def __getitem__(self, idx):
        return (torch.from_numpy(self.features[idx]), torch.from_numpy(self.
→labels[idx]))
```

6.2 NSLKDDFeatures

```
[ ]: class NSLKDDFeatures(Dataset):

    def __init__(self, features):
        self.features = np.asarray(features)

    def __len__(self):
        return self.features.shape[0]

    def __getitem__(self, idx):
        return torch.from_numpy(self.features[idx])
```

6.3 NSLKDDSiameseNetDataSet

```
[ ]: class NSLKDDSiameseNetDataSet(Dataset):

    def __init__(self, partitioned_data, indices, labels, nsl_kdd_features):
        self.partitioned_data = partitioned_data
        self.indices = indices
        self.labels = labels
        self.nsl_kdd_features = nsl_kdd_features

    def __len__(self):
        return len(self.indices)

    def __getitem__(self, idx):
```

```

        idx_tuple = self.indices[idx]
        if idx_tuple[0] == idx_tuple[1]:
            ret_list = self.partitioned_data[idx_tuple[0]] [
                np.random.choice(
                    self.partitioned_data[idx_tuple[0]].shape[0], 2,
→replace=False
                )
            ]
        else:
            ret_list = list(map(
                lambda x: self.partitioned_data[x] [
                    np.random.choice(self.partitioned_data[x].shape[0], 1)[0]
                ],
                idx_tuple)
            )

        return (
            torch.from_numpy(ret_list[0]),
            torch.from_numpy(ret_list[1]),
            torch.tensor(self.labels[idx], dtype=torch.int64)
        )

```

6.4 Training Data

```

[ ]: train_features_pd = pd.read_csv(os.path.join(CSV_DIR, 'training',
→'train_features_pca.csv'), header=None)
train_labels_multiclass_pd = pd.read_csv(os.path.join(CSV_DIR, 'training',
→'train_labels_multiclass.csv'), header=None)
train_labels_binary_pd = pd.read_csv(os.path.join(CSV_DIR, 'training',
→'train_labels_binary.csv'), header=None)

```

```

[ ]: train_features_pd.shape

```

```

[ ]: train_multiclass_ds = NSLKDDDataset(train_features, train_labels_multiclass)
train_binary_ds = NSLKDDDataset(train_features, train_labels_binary)

```

```

[ ]: train_features_ds = NSLKDDFeatures(train_features)

```

```

[ ]: train_multiclass_dl = DataLoader(
    train_multiclass_ds,
    batch_size=BATCH_SIZE,
    shuffle=True
)

train_binary_dl = DataLoader(

```

```

        train_binary_ds,
        batch_size=BATCH_SIZE,
        shuffle=True
    )

    train_features_dl = DataLoader(
        train_features_ds,
        batch_size=BATCH_SIZE,
        shuffle=True
    )

```

6.5 Testing Data

```

[ ]: test_features_pd = pd.read_csv(os.path.join(CSV_DIR, 'testing',
    ↳ 'test_features_pca.csv'), header=None)
test_labels_multiclass_pd = pd.read_csv(os.path.join(CSV_DIR, 'testing',
    ↳ 'test_labels_multiclass.csv'), header=None)
test_labels_binary_pd = pd.read_csv(os.path.join(CSV_DIR, 'testing',
    ↳ 'test_labels_binary.csv'), header=None)

```

```

[ ]: test_features_pd.shape

```

```

[ ]: test_multiclass_ds = NSLKDDDataset(test_features, test_labels_multiclass)
test_binary_ds = NSLKDDDataset(test_features, test_labels_binary)

```

```

[ ]: test_features_ds = NSLKDDFeatures(test_features)

```

```

[ ]: test_multiclass_dl = DataLoader(
    test_multiclass_ds,
    batch_size=BATCH_SIZE,
    shuffle=True
)

test_binary_dl = DataLoader(
    test_binary_ds,
    batch_size=BATCH_SIZE,
    shuffle=True
)

test_features_dl = DataLoader(
    test_features_ds,
    batch_size=BATCH_SIZE,
    shuffle=True
)

```

6.6 NumPy Arrays

```
[ ]: train_features_np = np.asarray(train_features_pd)
test_features_np = np.asarray(test_features_pd)

train_labels_binary_np = np.asarray(train_labels_binary_pd)
test_labels_binary_np = np.asarray(test_labels_binary_pd)

train_labels_multiclass_np = np.asarray(train_labels_multiclass_pd)
test_labels_multiclass_np = np.asarray(test_labels_multiclass_pd)
```

7 Data Plots

7.1 Plot Pie Chart

```
[ ]: def plot_pie_chart(labels_one_hot, title, figure_size=(8,8), have_title=True,
    ↳log_scaling=False):
    labels = np.argmax(labels_one_hot, axis=1)
    names = ['Normal', 'DoS', 'Probe', 'R2L', 'U2R']
    _, sizes = np.unique(labels, return_counts=True)

    names_sizes_dict = dict(zip(names, sizes))
    sorted_names_sizes_dict = {k: v for k, v in sorted(
        names_sizes_dict.items(),
        key=lambda item: item[1],
        reverse=True
    )}

    names = list(sorted_names_sizes_dict.keys())
    sizes = list(sorted_names_sizes_dict.values())

    if log_scaling:
        sizes = np.log(sizes)

    fig = plt.figure(figsize=figure_size)
    ax = fig.add_subplot()

    theme = plt.get_cmap('gray')
    ax.set_prop_cycle('color', [theme(1. * i / len(sizes))
        for i in range(len(sizes))])

    patches, texts, autotexts = ax.pie(
        sizes,
        autopct='%1.1f%%',
```

```

        startangle=90
    )

    autotexts[0].set_color('w')
    autotexts[1].set_color('w')

    # for text in autotexts:
    #     text.set_fontsize(11)

    ax.axis('equal') # Equal aspect ratio ensures that pie is drawn as a circle.

    if have_title:
        plt.title(
            title,
            fontsize=LEGEND_TEXT_SIZE
        )

    plt.legend(
        names,
        fontsize=11
    )

    plt.savefig(title + '.png')

```

```

[ ]: plot_pie_chart(
    train_labels_multiclass_np,
    title='Class Distributions: Training Data (log scaled)',
    figure_size=(10,10),
    have_title=False,
    log_scaling=True
)

```

```

[ ]: plot_pie_chart(
    test_labels_multiclass_np,
    title='Class Distributions: Testing Data (log scaled)',
    figure_size=(10,10),
    have_title=False,
    log_scaling=True
)

```

8 Model

8.1 Autoencoder

```
[ ]: class Autoencoder(nn.Module):

    def __init__(self, in_features, out_features):
        super(Autoencoder, self).__init__()

        self.encoder = nn.Sequential(
            nn.Linear(in_features=in_features, out_features=out_features),
            nn.Sigmoid()
        )

        self.decoder = nn.Sequential(
            nn.Linear(in_features=out_features, out_features=in_features),
            nn.Sigmoid()
        )

    def forward(self, x):
        x = self.encoder(x)
        x = self.decoder(x)

        return x
```

8.2 AnomalyDetector

```
[ ]: class AnomalyDetector(nn.Module):

    def __init__(self):
        super(AnomalyDetector, self).__init__()

        self.activation = nn.LeakyReLU(inplace=True)

        self.model = nn.Sequential(
            # nn.BatchNorm1d(num_features=67),
            nn.Linear(in_features=67, out_features=100),
            self.activation,

            nn.BatchNorm1d(num_features=100),
            nn.Linear(in_features=100, out_features=120),
            self.activation,

            nn.BatchNorm1d(num_features=120),
            nn.Linear(in_features=120, out_features=60),
```



```

        self.activation,

        nn.BatchNorm1d(num_features=60),
        nn.Linear(in_features=60, out_features=30),
        self.activation,

        nn.BatchNorm1d(num_features=30),
        nn.Linear(in_features=30, out_features=10),
        self.activation,

        nn.BatchNorm1d(num_features=10),
        nn.Linear(in_features=10, out_features=2),
        nn.Softmax(dim=1)
    )

    def forward(self, x):
        x = self.model(x)

        return x

```

8.3 SignatureGenerator

```

[ ]: class SignatureGenerator(nn.Module):

    def __init__(self):
        super(SignatureGenerator, self).__init__()

        self.activation = nn.LeakyReLU(inplace=True)

        self.model = nn.Sequential(
            nn.Linear(in_features=67, out_features=100),
            self.activation,
            nn.BatchNorm1d(num_features=100),

            nn.Linear(in_features=100, out_features=80),
            self.activation,
            nn.BatchNorm1d(num_features=80),

            nn.Linear(in_features=80, out_features=60),
            self.activation,
            nn.BatchNorm1d(num_features=60),

            nn.Linear(in_features=60, out_features=40),
            self.activation,
            nn.BatchNorm1d(num_features=40),

```

```

        nn.Linear(in_features=40, out_features=20),
        self.activation,
        nn.BatchNorm1d(num_features=20),

        nn.Linear(in_features=20, out_features=10),
        nn.Sigmoid()
    )

    def forward(self, x):
        x = self.model(x)

        return x

```

8.4 L1Difference

```

[ ]: class L1Difference(nn.Module):

    def __init__(self):
        super(L1Difference, self).__init__()

    def forward(self, x1, x2):
        return abs(x1 - x2)

```

8.5 SiameseNet

```

[ ]: class SiameseNet(nn.Module):

    def __init__(self):
        super(SiameseNet, self).__init__()

        self.signature_gen = SignatureGenerator()
        self.l1_diff = L1Difference()
        self.fc = nn.Linear(in_features=10, out_features=2)

    def forward(self, x1, x2):
        x1_sig = self.signature_gen(x1)
        x2_sig = self.signature_gen(x2)

        l1_diff = self.l1_diff(x1_sig, x2_sig)

        out = torch.softmax(self.fc(l1_diff), dim=1)

        return out

```

8.6 MultiClassClassifier

```
[ ]: class MultiClassClassifier(nn.Module):

    def __init__(self):
        super(MultiClassClassifier, self).__init__()

        self.model = nn.Sequential(
            nn.Linear(in_features=67, out_features=120),
            nn.ReLU(inplace=True),
            nn.BatchNorm1d(num_features=120),

            nn.Linear(in_features=120, out_features=100),
            nn.ReLU(inplace=True),
            nn.BatchNorm1d(num_features=100),

            nn.Linear(in_features=100, out_features=80),
            nn.ReLU(inplace=True),
            nn.BatchNorm1d(num_features=80),

            nn.Linear(in_features=80, out_features=40),
            nn.ReLU(inplace=True),
            nn.BatchNorm1d(num_features=40),

            nn.Linear(in_features=40, out_features=20),
            nn.ReLU(inplace=True),
            nn.BatchNorm1d(num_features=20),

            nn.Linear(in_features=20, out_features=10),
            nn.ReLU(inplace=True),
            nn.BatchNorm1d(num_features=10),

            nn.Linear(in_features=10, out_features=5),
            nn.Softmax(dim=1)
        )

    def forward(self, x):
        return self.model(x)
```

9 Training Utilities

9.1 Network Loss

```
[ ]: def get_network_loss(net, dataloader):
    agg_loss = 0.0

    criterion = nn.CrossEntropyLoss()

    for i, data in enumerate(dataloader, 0):
        x, y = data
        y = torch.argmax(y, dim=1)

        x = x.to(device)
        y = y.to(device)

        outputs = net(x)
        loss = criterion(outputs, y)
        agg_loss += loss.item()

    return agg_loss / len(dataloader)
```

9.2 Train Network

```
[ ]: def train_network(net, trainloader, testloader, n_epochs, print_every=500):
    net.double()
    net.to(device)

    optimizer = optim.Adam(net.parameters())
    criterion = nn.CrossEntropyLoss()

    train_losses = []
    test_losses = []
    for epoch in range(n_epochs):
        running_loss = 0.0
        for i, data in enumerate(trainloader):
            x, y = data
            y = torch.argmax(y, dim=1)

            x = x.to(device)
            y = y.to(device)

            optimizer.zero_grad()
            outputs = net(x)
            loss = criterion(outputs, y)
```

```

        loss.backward()
        optimizer.step()

        running_loss += loss.item()
        if i % print_every == print_every-1:
            print('[%d, %5d] loss: %.6f' %
                  (epoch + 1, i + 1, running_loss / print_every))
            running_loss = 0.0

    net.eval()
    train_loss = get_network_loss(net, trainloader)
    test_loss = get_network_loss(net, testloader)
    net.train()

    train_losses.append(train_loss)
    test_losses.append(test_loss)

    print('Training Loss: %.6f' % (train_loss))
    print('Testing Loss: %.6f' % (test_loss))

    return train_losses, test_losses

```

9.3 Encoder Loss

```

[ ]: def get_encoder_loss(autoencoder, dataloader, criterion, autoencoder1=None,
    ↪autoencoder2=None):
    agg_loss = 0.0

    for i, data in enumerate(dataloader, 0):
        x = data
        x = x.to(device)

        if autoencoder1:
            x = autoencoder1.encoder(x)
        if autoencoder2:
            x = autoencoder2.encoder(x)

        outputs = autoencoder(x)
        loss = criterion(outputs, x)
        agg_loss += loss.item()

    return agg_loss / len(dataloader)

```

9.4 Train Encoder

```
[ ]: def train_encoder(autoencoder, trainloader, testloader, n_epochs,
    print_every=500, autoencoder1=None, autoencoder2=None):
    autoencoder.double()
    autoencoder.to(device)

    optimizer = optim.Adam(autoencoder.parameters())
    criterion = nn.MSELoss()

    train_losses = []
    test_losses = []
    for epoch in range(n_epochs):
        running_loss = 0.0
        for i, data in enumerate(trainloader, 0):
            x = data
            x = x.to(device)

            with torch.no_grad():
                if autoencoder1:
                    x = autoencoder1.encoder(x)
                if autoencoder2:
                    x = autoencoder2.encoder(x)

            optimizer.zero_grad()
            outputs = autoencoder(x)
            loss = criterion(outputs, x)
            loss.backward()
            optimizer.step()

            running_loss += loss.item()
            if i % print_every == print_every-1:
                print('[%d, %5d] loss: %.6f' %
                    (epoch + 1, i + 1, running_loss / print_every))
                running_loss = 0.0

        autoencoder.eval()
        train_loss = get_encoder_loss(autoencoder, trainloader, criterion,
    autoencoder1=autoencoder1, autoencoder2=autoencoder2)
        test_loss = get_encoder_loss(autoencoder, testloader, criterion,
    autoencoder1=autoencoder1, autoencoder2=autoencoder2)
        autoencoder.train()

        train_losses.append(train_loss)
        test_losses.append(test_loss)

    print('Training Loss: %.6f' % (train_loss))
```

```

        print('Testing Loss: %.6f' % (test_loss))

    return train_losses, test_losses

```

9.5 SiameseNet Loss

```

[ ]: def get_siamese_net_loss(siamese_net, dataloader, val_batches):
    agg_loss = 0.0
    criterion = nn.CrossEntropyLoss()

    for i, data in enumerate(dataloader, 0):
        x1, x2, y = data

        x1 = x1.to(device)
        x2 = x2.to(device)
        y = y.to(device)

        outputs = siamese_net(x1.double(), x2.double())
        loss = criterion(outputs, y)
        agg_loss += loss.item()

        if i == val_batches-1:
            break

    return agg_loss / val_batches

```

9.6 Train SiameseNet

```

[ ]: def train_siamese_net(siamese_net, trainloader, testloader, n_epochs,
    print_every=500, val_batches=1):
    siamese_net.double()
    siamese_net.to(device)

    optimizer = optim.Adam(siamese_net.parameters())
    criterion = nn.CrossEntropyLoss()

    train_losses = []
    test_losses = []
    for epoch in range(n_epochs):
        running_loss = 0.0
        for i, data in enumerate(trainloader):
            x1, x2, y = data

            x1 = x1.to(device)

```

```

        x2 = x2.to(device)
        y = y.to(device)

        optimizer.zero_grad()
        outputs = siamese_net(x1.double(), x2.double())
        outputs = torch.squeeze(outputs)
        loss = criterion(outputs, y)
        loss.backward()
        optimizer.step()

        running_loss += loss.item()
        if i % print_every == print_every-1:
            print('%d, %5d loss: %.6f' %
                  (epoch + 1, i + 1, running_loss / print_every))
            running_loss = 0.0

    siamese_net.eval()
    train_loss = get_siamese_net_loss(siamese_net, trainloader, val_batches)
    test_loss = get_siamese_net_loss(siamese_net, testloader, val_batches)
    siamese_net.train()

    train_losses.append(train_loss)
    test_losses.append(test_loss)

    print('Training Loss: %.6f' % (train_loss))
    print('Testing Loss: %.6f' % (test_loss))

    return train_losses, test_losses

```

9.7 Partition Dataset

```

[ ]: def partition_dataset(data, labels):
    unique_labels = np.unique(labels)
    final_data = [None] * len(unique_labels)
    n_samples = len(labels)
    for i in range(n_samples):
        curr_class = labels[i]
        if final_data[curr_class]:
            final_data[curr_class].append(data[i])
        else:
            final_data[curr_class] = [data[i]]

    return list(map(lambda x: np.asarray(x), final_data))

```


9.8 Plotting Losses

```
[ ]: def plot_losses(train_losses, test_losses, title, xlabel, ylabel, start_idx=1,
    →interval_length=1, figure_size=(14,14)):
    train_losses_aggregated = []
    test_losses_aggregated = []
    n_epochs = len(train_losses)

    start_idx -= 1

    i = start_idx
    while i < n_epochs:
        curr_train_sum = sum(train_losses[i : min(n_epochs, i +
    →interval_length)])
        curr_test_sum = sum(test_losses[i : min(n_epochs, i + interval_length)])

        train_losses_aggregated.append(curr_train_sum / interval_length)
        test_losses_aggregated.append(curr_test_sum / interval_length)

        i += interval_length

    n = len(train_losses_aggregated)
    x = list(map(lambda x: (x * interval_length) + start_idx, np.arange(1, n+1)))

    plt.figure(figsize=figure_size)
    plt.plot(x, train_losses_aggregated, marker='o', linestyle='dashed')
    plt.plot(x, test_losses_aggregated, marker='o', linestyle='dashed')
    plt.legend(['Training Loss', 'Testing Loss'])
    plt.xlabel(xlabel)
    plt.ylabel(ylabel)
    plt.title(title)
    plt.grid()
    plt.xticks(x)

    plt.savefig(title + '.png')
    plt.show()
```

9.9 Saving and Loading Models

```
[ ]: def save_state_dict(model, filename):
    torch.save(model.state_dict(), filename)
```

```
[ ]: def load_state_dict(model, state_dict_path):
    state_dict = torch.load(state_dict_path)
    model.load_state_dict(state_dict)
```

```
model.double()
model.to(device)
```

10 Training

10.1 Autoencoders

10.1.1 Label Encoding

Autoencoder 1

```
[ ]: autoencoder_label_1 = Autoencoder(in_features=40, out_features=20)
    print(autoencoder_label_1)

[ ]: train_losses_encoder_label_1, test_losses_encoder_label_1 = train_encoder(
    autoencoder_label_1,
    trainloader=train_label_loader,
    testloader=test_label_loader,
    n_epochs=30
)

[ ]: plot_losses(
    train_losses_encoder_label_1,
    test_losses_encoder_label_1,
    # start_idx=11,
    title='autoencoder_label_1',
    xlabel='Epochs',
    ylabel='Losses',
    interval_length=1,
    figure_size=(14,8)
)

[ ]: save_state_dict(autoencoder_label_1, 'autoencoder_label_1.pth')
```

Autoencoder 2

```
[ ]: autoencoder_label_2 = Autoencoder(in_features=20, out_features=10)
    print(autoencoder_label_2)

[ ]: train_losses_encoder_label_2, test_losses_encoder_label_2 = train_encoder(
    autoencoder_label_2,
    autoencoder1=autoencoder_label_1,
    trainloader=train_label_loader,
    testloader=test_label_loader,
    n_epochs=30
)
```

```
[ ]: plot_losses(
    train_losses_encoder_label_2,
    test_losses_encoder_label_2,
    # start_idx=11,
    title='autoencoder_label_2',
    xlabel='Epochs',
    ylabel='Losses',
    interval_length=1,
    figure_size=(14,8)
)
```

```
[ ]: save_state_dict(autoencoder_label_2, 'autoencoder_label_2.pth')
```

10.2 Networks

```
[ ]: LOAD_DIR = '/content/drive/My Drive/Colab Notebooks/Intrusion Detection/
    ↳saved_models'
```

10.2.1 Loading Autoencoder 1

```
[ ]: autoencoder_label_1 = Autoencoder(in_features=40, out_features=20)
    load_state_dict(autoencoder_label_1, os.path.join(LOAD_DIR, 'autoencoder_label_1.
    ↳pth'))
```

10.2.2 Loading Autoencoder 2

```
[ ]: autoencoder_label_2 = Autoencoder(in_features=20, out_features=10)
    load_state_dict(autoencoder_label_2, os.path.join(LOAD_DIR, 'autoencoder_label_2.
    ↳pth'))
```

10.2.3 AnomalyDetector

```
[ ]: anomaly_detector = AnomalyDetector()

    print(anomaly_detector)
```

```
[ ]: train_losses_anomaly_detector, test_losses_anomaly_detector = train_network(
    anomaly_detector,
    trainloader=train_binary_dl,
    testloader=test_binary_dl,
    n_epochs=50,
```

```
    print_every=100
)
```

```
[ ]: plot_losses(
    train_losses_network_label,
    test_losses_network_label,
    # start_idx=37,
    title='network_label_double_encoder',
    xlabel='Epochs',
    ylabel='Losses',
    interval_length=1,
    figure_size=(14,8))
```

10.3 Data Distributions

10.3.1 PCA

```
[ ]: pca = PCA(n_components=3)
pca.fit(train_features_np)
```

```
[ ]: train_features_3dims = pca.transform(train_features_np)
test_features_3dims = pca.transform(test_features_np)
```

10.3.2 Groups

```
[ ]: groups_multiclass = [
    'Normal',
    'DoS',
    'Probe',
    'R2L',
    'U2R'
]

groups_binary = [
    'Normal',
    'Malicious'
]

groups_malicious_classes = [
    'DoS',
    'Probe',
    'R2L',
    'U2R'
]
```

10.3.3 Plot Points

```
[ ]: def plot_points(
    features,
    labels,
    title,
    groups,
    class_idx=None,
    one_hot_labels=True,
    figure_size=(8,8)
):

    fig = plt.figure(figsize=figure_size)
    ax = fig.add_subplot(111, projection='3d')

    if one_hot_labels:
        labels_indices = np.asarray(list(map(lambda x: np.argmax(x), labels)))
    else:
        labels_indices = labels

    data = []
    if class_idx:
        i = class_idx
        curr_data_indices = np.where(labels_indices == i)[0]
        curr_data = features[curr_data_indices]
        data.append(curr_data)
    else:
        for i in range(len(groups)):
            curr_data_indices = np.where(labels_indices == i)[0]
            curr_data = features[curr_data_indices]
            data.append(curr_data)

    all_colors = (
        'grey',
        'orange',
        'cyan',
        'green',
        'maroon',
        'chocolate',
        'lightgreen',
        'dodgerblue',
        'darkslategray',
        'lightseagreen',
        'mediumspringgreen',
        'rebeccapurple',
        'hotpink',
        'indigo',
```

```

        'midnightblue',
        'gold',
        'black',
        'crimson',
        'darkkhaki',
        'aqua'
    )

    colors = sample(all_colors, len(groups))

    for features, color, group in zip(data, colors, groups):
        x = list(map(lambda x: x[0], features))
        y = list(map(lambda x: x[1], features))
        z = list(map(lambda x: x[2], features))
        ax.scatter(
            x, y, z,
            c=color,
            label=group,
        )

    plt.legend()
    plt.title(
        title,
        fontsize=15,
        loc='left'
    )

    plt.savefig(title + '.png')

```

10.3.4 Training Data Plots

Binary Class Distribution

```

[ ]: plot_points(
    train_features_3dims,
    train_labels_binary_np,
    groups=groups_binary,
    title='Binary Class Distribution',
    figure_size=(14,14)
)

```

Multi Class Distribution

```

[ ]: plot_points(
    train_features_3dims,
    train_labels_multiclass_np,

```

```
groups=groups_multiclass,  
title='Multi Class Distribution',  
figure_size=(14,14)  
)
```

Malicious Data Point Distributions

DoS

```
[ ]: plot_points(  
    train_features_3dims,  
    train_labels_multiclass_np,  
    class_idx=1,  
    groups=['DoS'],  
    title='DoS',  
    figure_size=(14,14)  
)
```

Probe

```
[ ]: plot_points(  
    train_features_3dims,  
    train_labels_multiclass_np,  
    class_idx=2,  
    groups=['Probe'],  
    title='Probe',  
    figure_size=(14,14)  
)
```

R2L

```
[ ]: plot_points(  
    train_features_3dims,  
    train_labels_multiclass_np,  
    class_idx=3,  
    groups=['R2L'],  
    title='R2L',  
    figure_size=(14,14)  
)
```

U2R

```
[ ]: plot_points(  
    train_features_3dims,  
    train_labels_multiclass_np,  
    class_idx=4,
```

```

groups=['U2R'],
title='U2R',
figure_size=(14,14)
)

```

10.4 Anomaly Detection

10.4.1 Utilities

Plot Confusion Matrix

```

[ ]: def plot_confusion_matrix(
    cm, classes,
    normalize=False,
    title='Confusion Matrix',
    cmap=plt.cm.Blues,
    print_cm=False,
    save=False
):
    """
    This function prints and plots the confusion matrix.
    Normalization can be applied by setting `normalize=True`.
    """

    import matplotlib.pyplot as plt
    import numpy as np
    import itertools

    if normalize:
        cm = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]
        print("Normalized confusion matrix")
    else:
        print('Confusion matrix, without normalization')

    if print_cm:
        print(cm)

    plt.imshow(cm, interpolation='nearest', cmap=cmap)
    # plt.title(title)
    plt.colorbar()
    tick_marks = np.arange(len(classes))
    plt.xticks(tick_marks, classes, rotation=45)
    plt.yticks(tick_marks, classes)

    fmt = '.2f' if normalize else 'd'
    thresh = cm.max() / 2.
    for i, j in itertools.product(range(cm.shape[0]), range(cm.shape[1])):

```



```

plt.text(j, i, format(cm[i, j], fmt),
         horizontalalignment="center",
         color="white" if cm[i, j] > thresh else "black")

plt.ylabel('True')
plt.xlabel('Predicted')
# plt.xticks([])
# plt.yticks([])
plt.grid(False)

if save:
    plt.savefig(title + '.png')

```

One Hot to Label Encoding

```

[ ]: def to_label(arr):
      return np.argmax(arr, axis=1)

```

Get Predictions

```

[ ]: def get_predictions(estimator, features, labels):
      return to_label(labels), estimator.predict(features), estimator.
      ↪predict_proba(features)

```

Get Scores

```

[ ]: def get_scores(y_true, y_pred):
      accuracy = accuracy_score(y_true, y_pred)
      precision = precision_score(y_true, y_pred)
      recall = recall_score(y_true, y_pred)
      f1 = f1_score(y_true, y_pred)

      return {
          'accuracy': accuracy,
          'precision': precision,
          'recall': recall,
          'f1': f1
      }

```

Get Adjusted Classes

```

[ ]: def get_adjusted_classes(y_probs, thresh=0.5):
      return [int(y_prob >= thresh) for y_prob in y_probs]

```

Plot ROC Curve

```
[ ]: def plot_roc_curve(fpr, tpr, title, label=None, figure_size=(8,8)):
    """
    The ROC curve, modified from
    Hands-On Machine learning with Scikit-Learn and TensorFlow; p.91
    """

    plt.figure(figsize=figure_size)
    plt.title('ROC Curve')
    plt.plot(fpr, tpr, linewidth=2, label=label)
    plt.plot([0, 1], [0, 1], 'k--')
    plt.axis([-0.005, 1, 0, 1.005])
    plt.xticks(np.arange(0,1, 0.05), rotation=90)
    plt.xlabel("False Positive Rate")
    plt.ylabel("True Positive Rate (Recall)")
    plt.legend(loc='best')

    plt.savefig(title + '.png')
```

Plot Precision Vs Recall

```
[ ]: def plot_precision_recall_curve(p, r, y_scores, y_test, thresholds, title, t=0.
    ↪5, figure_size=(8,8)):
    """
    Plots the precision recall curve and shows the current value for each
    by identifying the classifier's threshold (t).
    """

    # generate new class predictions based on the adjusted_classes
    # function above and view the resulting confusion matrix.
    y_pred_adj = get_adjusted_classes(y_scores, t)

    # plot the curve
    plt.figure(figsize=figure_size)
    plt.title(title)
    plt.step(r, p, color='b', alpha=0.2,
            where='post')
    plt.fill_between(r, p, step='post', alpha=0.2,
                    color='b')
    plt.ylim([0.5, 1.01]);
    plt.xlim([0.5, 1.01]);
    plt.xlabel('Recall');
    plt.ylabel('Precision');

    # plot the current threshold on the line
    close_default_clf = np.argmin(np.abs(thresholds - t))
    plt.plot(r[close_default_clf] - 0.01, p[close_default_clf], '.', c='k',
            markersize=15)
    plt.text(r[close_default_clf], p[close_default_clf], 'Current Threshold')
```

```
plt.savefig(title + '.png')
```

Plot Precision and Recall Vs Threshold

```
[ ]: def plot_precision_recall_vs_threshold(precisions, recalls, thresholds, title,
      →figure_size=(8,8)):
    """
    Modified from:
    Hands-On Machine learning with Scikit-Learn
    and TensorFlow; p.89
    """
    idx = np.argmin(abs(precisions - recalls))
    plt.figure(figsize=figure_size)
    plt.title(title)
    plt.plot(thresholds, precisions[:-1], "k--", label="Precision")
    plt.plot(thresholds, recalls[:-1], "k-", label="Recall")
    plt.plot(thresholds[idx], precisions[idx], 'ro')
    plt.text(thresholds[idx] + 0.01, precisions[idx], f'Threshold = {
      →thresholds[idx]}')
    plt.ylabel("Score")
    plt.xlabel("Decision Threshold")
    plt.legend(loc='best')

    plt.savefig(title + '.png')

    return thresholds[idx]
```

10.4.2 Clustering Normal Points

```
[ ]: plot_points(
    train_features_3dims,
    train_labels_multiclass_np,
    class_idx=0,
    groups=['Normal'],
    title='Normal',
    figure_size=(14,14)
)
```

```
[ ]: train_labels = np.asarray(list(map(lambda x: np.argmax(x),
      →train_labels_binary_np)))
train_indices_normal = np.where(train_labels == 0)[0]
train_features_normal = train_features_np[train_indices_normal]
```

```
test_labels = np.asarray(list(map(lambda x: np.argmax(x),
    ↳test_labels_binary_np)))
test_indices_normal = np.where(test_labels == 0)[0]
test_features_normal = test_features_np[test_indices_normal]
```

```
[ ]: train_features_normal.shape
```

```
[ ]: test_features_normal.shape
```

K-Means

Finding Optimal Clusters using K-Elbow Visualizer

```
[ ]: model = cluster.KMeans(random_state=42)
visualizer = KElbowVisualizer(
    model,
    k=(2,11),
    metric='calinski_harabasz'
)

visualizer.fit(train_features_normal)
```

Fitting

```
[ ]: N_CLUSTERS = 4
```

```
[ ]: kmeans_normal = cluster.KMeans(n_clusters=N_CLUSTERS, random_state=42)
kmeans_normal.fit(train_features_normal)
```

```
[ ]: kmeans_normal.inertia_
```

```
[ ]: train_labels_normal = kmeans_normal.predict(train_features_normal)
```

```
[ ]: plot_points(
    features=pca.transform(train_features_normal),
    labels=train_labels_normal,
    groups=['Cluster ' + str(i+1) for i in range(N_CLUSTERS)],
    one_hot_labels=False,
    title='K-Means Clustering on Normal Data Points - Training',
    figure_size=(14,14)
)
```

```
[ ]: test_labels_normal = kmeans_normal.predict(test_features_normal)
```

```
[ ]: plot_points(
    features=pca.transform(test_features_normal),
    labels=test_labels_normal,
    groups=['Cluster ' + str(i+1) for i in range(N_CLUSTERS)],
    one_hot_labels=False,
    title='K-Means Clustering on Normal Data Points - Testing',
    figure_size=(14,14)
)

[ ]: joblib.dump(kmeans_normal, 'kmeans_normal.joblib')
```

Partitioning Dataset

```
[ ]: partitioned_training_data_normal = partition_dataset(train_features_normal,
    →train_labels_normal)

[ ]: for data in partitioned_training_data_normal:
    print(data.shape)

[ ]: partitioned_testing_data_normal = partition_dataset(test_features_normal,
    →test_labels_normal)

[ ]: for data in partitioned_testing_data_normal:
    print(data.shape)
```

SiameseNet Generating Data for NSLKDDSiameseNetDataset Instance

```
[ ]: TRAIN_MULT_FACTOR = 100000
    TEST_MULT_FACTOR = 20000

[ ]: indices = []
    for i in range(N_CLUSTERS):
        for j in range(i, N_CLUSTERS):
            indices.append((i,j))

    train_indices_final = indices * TRAIN_MULT_FACTOR
    shuffle(train_indices_final)
    test_indices_final = indices * TEST_MULT_FACTOR
    shuffle(test_indices_final)

[ ]: len(train_indices_final), len(test_indices_final)

[ ]: train_labels = list(map(lambda x: int(x[0] == x[1]), train_indices_final))
    test_labels = list(map(lambda x: int(x[0] == x[1]), test_indices_final))
```

```
[ ]: len(train_labels), len(test_labels)
```

NSLKDDSiameseNetDataset Instance

```
[ ]: train_data_siamese_net = NSLKDDSiameseNetDataSet(  
    partitioned_training_data_normal,  
    train_indices_final,  
    np.asarray(train_labels),  
    train_features  
)
```

```
[ ]: trainloader_siamese_net = DataLoader(  
    train_data_siamese_net,  
    batch_size=256,  
    shuffle=True  
)
```

```
[ ]: test_data_siamese_net = NSLKDDSiameseNetDataSet(  
    partitioned_testing_data_normal,  
    test_indices_final,  
    np.asarray(test_labels),  
    train_features  
)
```

```
[ ]: testloader_siamese_net = DataLoader(  
    test_data_siamese_net,  
    batch_size=256,  
    shuffle=True  
)
```

Training Signatures for Each Cluster

```
[ ]: siamese_net = SiameseNet()  
     print(siamese_net)
```

```
[ ]: train_losses_siamese_net, test_losses_siamese_net = train_siamese_net(  
    siamese_net,  
    trainloader=trainloader_siamese_net,  
    testloader=testloader_siamese_net,  
    n_epochs=50  
)
```

```
[ ]:
```

Z-Score

```
[ ]: curr_cluster = partitioned_training_data_normal[0]
```

```
[ ]: curr_cluster.shape
[ ]: z_scores = zscore(curr_cluster)
[ ]: z_scores.shape
[ ]: sns.kdeplot(z_scores[:,5])
```

10.4.3 Random Forest Classifier

Grid Search

Recall

```
[ ]: rfc = RandomForestClassifier(
    class_weight='balanced',
    max_depth=58,
    min_samples_leaf=2
)

param_grid = {
    'n_estimators': np.arange(100, 120),
}

rfc_grid_search_recall = GridSearchCV(
    estimator=rfc,
    param_grid=param_grid,
    scoring=make_scorer(recall_score),
    verbose=2,
    cv=5,
    n_jobs=5
)

[ ]: rfc_grid_search_recall.fit(train_features_np, to_label(train_labels_binary_np))
[ ]: rfc_grid_search_recall.best_score_
[ ]: rfc_grid_search_recall.best_params_
[ ]: joblib.dump(rfc_grid_search_recall, 'rfc_grid_search_recall.joblib')
```

Model

```
[ ]: rfc_grid_search_recall = joblib.load(os.path.join(MODELS_DIR,
    → 'rfc_grid_search_recall.joblib'))
```

```
[ ]: rfc = rfc_grid_search_recall.best_estimator_
```

```
y_true, y_pred, y_probs = get_predictions(  
    rfc,  
    test_features_np,  
    test_labels_binary_np  
)
```

```
[ ]: cm = confusion_matrix(  
    y_pred=y_pred,  
    y_true=y_true  
)
```

```
[ ]: plot_confusion_matrix(  
    cm,  
    classes=[0,1],  
    title='Random Forest Classifier with Optimized Recall',  
    save=True,  
    normalize=True,  
    cmap=plt.cm.Greys  
)
```

```
[ ]: p, r, thresholds = precision_recall_curve(y_true, y_probs[:,1])
```

```
[ ]: best_threshold_rfc = plot_precision_recall_vs_threshold(  
    p, r, thresholds,  
    title='Precision and Recall Scores as a Function of the Decision Threshold',  
)
```

```
[ ]: plot_precision_recall_curve(  
    p, r, y_probs[:,1], y_true, thresholds,  
    title='Precision Recall Curve',  
    t=best_threshold_rfc,  
)
```

```
[ ]: fpr, tpr, auc_thresholds = roc_curve(y_true, y_probs[:,1])
```

```
[ ]: auc(fpr, tpr)
```

```
[ ]: plot_roc_curve(  
    fpr, tpr,  
    title='ROC Curve',  
    label='Recall Optimized'  
)
```

```
[ ]: y_pred_adjusted = get_adjusted_classes(y_probs[:,1], thresh=best_threshold_rfc)
```



```
[ ]: cm = confusion_matrix(
    y_pred=y_pred_adjusted,
    y_true=y_true
)
```

```
[ ]: plot_confusion_matrix(
    cm,
    classes=[0,1],
    title='Random Forest Classifier with Optimized Recall',
    save=True,
    normalize=True,
    cmap=plt.cm.Greys
)
```

10.4.4 AdaBoost Classifier

Grid Search

Recall

```
[ ]: adaboost_clf = AdaBoostClassifier()

param_grid = {
    'n_estimators': np.arange(50, 70)
}

adaboost_clf_grid_search_recall = GridSearchCV(
    estimator=adaboost_clf,
    param_grid=param_grid,
    scoring=make_scorer(recall_score),
    verbose=2,
    cv=5,
    n_jobs=5
)
```

```
[ ]: adaboost_clf_grid_search_recall.fit(train_features_np,
    →to_label(train_labels_binary_np))
```

```
[ ]: joblib.dump(adaboost_clf_grid_search_recall, 'adaboost_clf_grid_search_recall.
    →joblib')
```

Model

```
[ ]: adaboost_clf_grid_search_recall = joblib.load(os.path.join(MODELS_DIR,
    →'adaboost_clf_grid_search_recall.joblib'))
```

```
[ ]: adaboost_clf = adaboost_clf_grid_search_recall.best_estimator_  
  
y_true, y_pred, y_probs = get_predictions(  
    adaboost_clf,  
    test_features_np,  
    test_labels_binary_np  
)
```

```
[ ]: cm = confusion_matrix(  
    y_pred=y_pred,  
    y_true=y_true  
)
```

```
[ ]: plot_confusion_matrix(  
    cm,  
    classes=[0,1],  
    title='AdaBoost Classifier with Optimized Recall',  
    save=True,  
    normalize=True,  
    cmap=plt.cm.Greys  
)
```

```
[ ]: p, r, thresholds = precision_recall_curve(y_true, y_probs[:,1])
```

```
[ ]: best_threshold = plot_precision_recall_vs_threshold(  
    p, r, thresholds,  
    title='Precision and Recall Scores as a Function of the Decision Threshold',  
)
```

```
[ ]: plot_precision_recall_curve(  
    p, r, y_probs[:,1], y_true, thresholds,  
    title='Precision Recall Curve',  
    t=best_threshold,  
)
```

```
[ ]: fpr, tpr, auc_thresholds = roc_curve(y_true, y_probs[:,1])
```

```
[ ]: auc(fpr, tpr)
```

```
[ ]: plot_roc_curve(  
    fpr, tpr,  
    title='ROC Curve',  
    label='Recall Optimized'  
)
```

```
[ ]: y_pred_adjusted = get_adjusted_classes(y_probs[:,1], thresh=best_threshold)
```

```
[ ]: cm = confusion_matrix(  
    y_pred=y_pred_adjusted,  
    y_true=y_true  
)
```

```
[ ]: plot_confusion_matrix(  
    cm,  
    classes=[0,1],  
    title='AdaBoost Classifier with Optimized Recall',  
    save=True,  
    normalize=True,  
    cmap=plt.cm.Greys  
)
```

```
[ ]:
```

10.5 MultiClassClassifier

Utilities

```
[ ]: def print_counts(labels):  
    labels = np.argmax(labels, axis=1)  
    for i in range(5):  
        print(len(np.where(labels == i)[0]))
```

```
[ ]: print_counts(test_labels_multiclass_np)
```

Model

```
[ ]: over_sampler = SVMSMOTE()
```

```
[ ]: x_train, y_train = over_sampler.fit_resample(  
    train_features_np, np.argmax(  
        train_labels_multiclass_np,  
        axis=1  
    )  
)
```

```
[ ]: x_train, y_train = shuffle(x_train, y_train)
```

```
[ ]: Counter(y_train)
```

```
[ ]: x_test, y_test = test_features_np, test_labels_multiclass_np
```

```
[ ]: x_train.shape, x_test.shape, y_train.shape, y_test.shape
```

```
[ ]: lstm_output_size = 70
```

```
[ ]: cnn = Sequential()
cnn.add(Convolution1D(64, 3,
    ↳border_mode="same",activation="relu",input_shape=(67, 1)))
cnn.add(Convolution1D(64, 3, border_mode="same", activation="relu"))
cnn.add(MaxPooling1D(pool_length=(2)))
cnn.add(Convolution1D(128, 3, border_mode="same", activation="relu"))
cnn.add(Convolution1D(128, 3, border_mode="same", activation="relu"))
cnn.add(MaxPooling1D(pool_length=(2)))
cnn.add(LSTM(lstm_output_size))
cnn.add(Dropout(0.1))
cnn.add(Dense(5, activation="softmax"))
```

```
[ ]: cnn.summary()
```

```
[ ]: cnn.compile(loss='categorical_crossentropy', optimizer='adam',
    ↳metrics=['accuracy'])
```

```
[ ]: history_1 = cnn.fit(
    np.expand_dims(x_train, axis=2), to_categorical(
        y_train,
        num_classes=5
    ),
    batch_size=128,
    epochs=5,
    validation_data=(
        np.expand_dims(x_test, axis=2), y_test
    )
)
```

```
[ ]: all_losses = []
```

```
[ ]: for loss in history_1.history['loss']:
    all_losses.append(loss)
```

```
[ ]: all_acc = []
```

```
[ ]: for acc in history_1.history['accuracy']:
    all_acc.append(acc)
```

```
[ ]: len(all_losses), len(all_acc)
```

```
[ ]: plt.figure(figsize=(8,8))

plt.plot(
    np.linspace(1, 10, 10) * 10,
```

```

        all_losses,
        'k--',
        marker='D'
    )

    plt.xlabel('Epochs')
    plt.ylabel('Loss')

    plt.savefig('Loss.png')

```

```

[ ]: plt.figure(figsize=(8,8))

plt.plot(
    np.linspace(1, 10, 10) * 10,
    all_acc,
    'k--',
    marker='D'
)

plt.xlabel('Epochs')
plt.ylabel('Accuracy')

plt.savefig('Accuracy.png')

```

```

[ ]: cnn.save('conv_model.h5')

```

```

[ ]: import pickle as pk

with open('history.pkl', 'wb') as f:
    pk.dump(history, f)

```

```

[ ]: import pickle as pk

with open('history_1.pkl', 'wb') as f:
    pk.dump(history_1, f)

```

```

[ ]: plot_model(cnn, show_layer_names=False, show_shapes=True, t)

```

```

[ ]: y_pred_rfc = rfc.predict(test_features_np)

```

```

[ ]: normal_samples_indices = np.where(y_pred_rfc == 0)[0]
x_train_normal = test_features_np[normal_samples_indices]
y_train_multiclass_normal = test_labels_multiclass_np[normal_samples_indices]

```

```

[ ]: y_train_multiclass_normal

```

```

[ ]: x_train_normal.shape, y_train_multiclass_normal.shape

```

```
[ ]: malicious_samples_indices = np.where(y_pred_rfc == 1)[0]
x_train_malicious = test_features_np[malicious_samples_indices]
y_train_multiclass_malicious =
    ↳test_labels_multiclass_np[malicious_samples_indices]

[ ]: x_train_malicious.shape, y_train_multiclass_malicious.shape

[ ]: y_true, y_pred, _ = get_predictions(
    cnn,
    np.expand_dims(x_train_malicious, axis=2),
    y_train_multiclass_malicious
)

[ ]: y_true_with_zeros = np.concatenate((y_true, np.argmax(y_train_multiclass_normal,
    ↳axis=1)))

[ ]: y_true_with_zeros.shape

[ ]: pred = np.argmax(y_pred, axis=1)

[ ]: pred_with_zeros = np.concatenate((pred, np.asarray([0] * 13920)))

[ ]: pred_with_zeros.shape

[ ]: pred_with_zeros

[ ]: cm = confusion_matrix(
    y_pred=pred_with_zeros,
    y_true=y_true_with_zeros
)

[ ]: plot_confusion_matrix(
    cm,
    classes=['Normal', 'DoS', 'Probe', 'U2R', 'R2L'],
    title='Final Confusion',
    save=True,
    normalize=True,
    cmap=plt.cm.Greys
)

[ ]:
```